

Joint Rc and RefCell Analysis in Rust

ALEXANDER CARMICHAEL, University of Colorado Boulder, USA

I present a combined static analyzer for Rust. Rc and RefCell are frequently used together in order to implement recursive data structures in Rust. My analyzer combines analyzing both Rc and RefCell in a single pass, which finds errors with the usage of both primitives. Since common runtime errors stemming from improper usage of RefCell are similar to that of Rc, this process has the potential to find errors with the usage of both, with less code, and a faster runtime than checking them individually.

ACM Reference Format:

Alexander Carmichael. 2025. Joint Rc and RefCell Analysis in Rust. 1, 1 (December 2025), 4 pages. <https://doi.org/10.1145/nmnnnn.nmnnnn>

1 Introduction

Rust is an incredibly popular statically typed programming language. Rust includes a built in borrow checker to guard against common memory errors, such as use-after-free. The borrow checker requires that every variable have a single "owner". Other data structures can "borrow" or "mutably borrow" a variable in order to read and write to it. A variable cannot be edited by its owner or mutably borrowed if at least one regular borrow is active, in order to prevent a data race. This rule makes it impossible to implement recursive data structures in Rust in the same manner as they would be done in a language such as C++. For that reason, Rust includes several primitives to get around the borrow checker's restrictions. They still have the same requirements as the borrow checker, but shift the verification from compile-time to runtime. In a data structure like a graph, multiple nodes can point to a child node. Rust's borrow checker makes that difficult. One node must be the "owner", and every other node that points to the child node must prove that it's lifetime is less than or equal to that of the owner. Rust provides the Rc primitive to address that requirement. Rc contains an immutable inner value, and can have many owners. Rc maintains an internal count of owners. When owners go out of scope, Rc decrements its internal counter. When the number of owners reaches zero, Rc deallocates its internal value and itself. This allows many structures to point to a single node without having to prove anything to the borrow checker. Rc has a restriction in that the internal data is immutable. Another primitive, RefCell, is commonly used to alleviate that. RefCell contains something mutable, and provides the .borrow() and .borrow_mut() methods to access it. These methods are dynamically checked, but obey the same rules as the borrow checker. A mutable borrow when a borrow is already active will crash the program.

1.1 Motivating Example

```
01 use std::rc::Rc;
02 use std::cell::RefCell;
03
```

Author's Contact Information: Alexander Carmichael, Alexander.Carmichael@colorado.edu, University of Colorado Boulder, Boulder, Colorado, USA.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM XXXX-XXXX/2025/12-ART

<https://doi.org/10.1145/nmnnnn.nmnnnn>

```

50  04 struct Node {
51    05   id: i32,
52    06   next: Option<Rc<RefCell<Node>>>,
53    07 }
54  08
55  09 impl Drop for Node {
56    11   fn drop(&mut self) {
57      12     println!("Node {} dropped.", self.id);
58    13   }
59  14 }
60  15
61  16 fn main() {
62    17   let a = Rc::new(RefCell::new(Node { id: 1, next: None }));
63  18
64  19   let b = Rc::new(RefCell::new(Node { id: 2, next: None }));
65  20
66  21   a.borrow_mut().next = Some(b.clone());
67  22   b.borrow_mut().next = Some(a.clone());
68  23
69  24   drop(a);
70  25   drop(b);
71  26 }
```

72 This program allocates two "nodes". It then makes the next node of each, the other. It then drops
 73 both nodes. What is interesting is that the "Node _ dropped." message will never print. That's
 74 because the nodes are never actually deallocated. The fact that each node points to the other means
 75 that each node's Rc will have a strong reference count of one. An Rc only deallocates its inside
 76 when its strong reference count reaches zero. That means these nodes will remain as long as this
 77 program is running, a memory leak.

```

78
79  01 fn main() {
80    02   let a = Rc::new(RefCell::new(Node { id: 1, next: None }));
81  03
82  04   let b = a.borrow();
83  05   let c = a.borrow_mut();
84  }
```

85 This second motivating example shows a potential error when using RefCell. The RefCell must
 86 abide by the normal borrowing rules. A mutable borrow is not allowed when there is at least one
 87 other borrow, as that could lead to a memory error such as use after free. Since RefCell is checked
 88 at runtime, this code passes Rust's borrow checker and compiles. When run, line 5 crashes the
 89 program as the requirement of no other borrows is not met.

1.2 Contribution

91 A dual-purpose static analyzer for Rust.

- 92 • An analyzer for Rc errors. It keeps an internal graph of all active variables and the variables
 93 that they point to. Upon the deallocation of a variable, it checks to see if the variable is still
 94 somehow accessible from the code. If it is not, it then checks the strong reference count to
 95 the variable. If it's greater than 0, then there is a memory leak, and the analyzer detects it.

- 99 • An analyzer for RefCell errors. Building on top of the first analyzer, it also keeps a record of
 100 the number of borrows made for each struct in the graph. If the combination of borrows
 101 made violates Rust's borrow checking rules, the analyzer flags it as a location where the
 102 program will crash.

103 The idea is to check both of these similar but different properties together. The benefit is that
 104 the shared work between the analyses can be factored out, and therefore the overall analysis is
 105 easier to write.

107 2 Semantics

$$\begin{array}{ll} \text{struct } s & ::= \{f_1 : s_1, f_2 : s_2, f_3 : s_3, \dots, f_n : s_n\} \\ \text{borrow } b & ::= b\{s\} \\ \text{mutable borrow } m & ::= m\{s\} \\ \text{variable } v & ::= s \mid b \mid m \\ \text{all vars } a & ::= \{v, v_2, v_3, \dots, v_n\} \end{array}$$

$$\frac{\boxed{a \vdash e : a}}{\begin{array}{c} \text{FIELDASSIGN} \\ a \vdash x \Downarrow s \quad a[v] \Downarrow s \\ \hline a \vdash s.f = x \Downarrow a\{v \rightarrow s\{f \rightarrow x\}\} \end{array}} \quad \frac{\text{DROPSTRUCT} \\ a[v] \Downarrow s}{a \vdash \text{drop}(v) \Downarrow a.\text{drop}(v)}$$

$$\frac{\text{VARASSIGNSTRUCT} \\ a \vdash x \Downarrow s \mid b \mid m}{a \vdash \text{let } v = x \Downarrow a\{v \rightarrow x\}}$$

$$\frac{\text{VARASSIGNBORROW} \\ \forall y \in a, a[y] \neq m\{x\}}{a \vdash \text{let } v = b\{x\} \Downarrow a\{v \rightarrow b\{x\}\}}$$

$$\frac{\text{VARASSIGNBORROWMUTABLE} \\ \forall y \in a, a[y] \neq b\{x\}, a[y] \neq m\{x\}}{a \vdash \text{let } v = b\{x\} \Downarrow a\{v \rightarrow b\{x\}\}}$$

125 3 Conclusion

126 The analyzer works. When given the first motivating example, it outputs "dropping b leaked
 127 memory as b still has remaining references". When given the second, "borrowing a failed because
 128 of an invalid borrow configuration". The actual Rust implementation is somewhat brittle, in that it
 129 doesn't comprehend enough of the Rust AST to traverse programs with irrelevant statements. It
 130 does not support Rust's more complicated methods/patterns for variable initialization, and it will
 131 therefore get confused/fail on more complicated syntax. A lot is left to be implemented. Recursive
 132 data structures are frequently initialized through loops, so support of loops would be a significant
 133 improvement. As for the actual goals of the project, this analyzer was somewhat of a success and
 134 somewhat of a failure. It's a success in that it does successfully find both types of issues, and there
 135 was a lot of similarity between the issues that was successfully factored out. The problem is that
 136 RefCell errors are rather simplistic by nature. It's difficult to come up with examples where misuse
 137 of RefCell results in a program crash that aren't just the programmer borrowing, forgetting they
 138 already have the value, and borrowing again to look at it without freeing or reusing the first. It's
 139 hard to make a more complicated error as that essentially requires "fighting" Rust's built in borrow
 140 checker, which already does much to make dangling references and other memory errors impossible.
 141 The result is that the Rc side of the analyzer is somewhat complex, and there are many ways to
 142 improve it and make it more useful. The RefCell side of the analyzer is simplistic and already
 143 does everything it can. Any further attempt to find issues with RefCell would be uninteresting, as
 144 those occur extremely rarely and require willful ignorance of Rust design patterns and the borrow
 145 checker's advice. So yes, while this analyzer does succeed in sharing some logic between the RefCell
 146 checker, it's still a bit of a mess.

148 and Rc analyses, it's rather limited, and there isn't any room for further sharing of logic. I still
149 think the idea has merit. RefCell in part has simple errors because it itself is a simple primitive.
150 Perhaps Rust's more complex primitives, meant for multi-threaded use, have the answer. Maybe a
151 combination analyzer for a multi-threaded ownership primitive, and a multi-threaded borrowing
152 primitive has the needed complexity to make a multi-purpose analyzer that works and has enough
153 overlap that something is gained by doing the analysis together.

154

155

156

157

158

159

160

161

162

163

164

165

166

167

168

169

170

171

172

173

174

175

176

177

178

179

180

181

182

183

184

185

186

187

188

189

190

191

192

193

194

195

196